Rajeev Ranjan

Senior Data Engineer | Azure | Databricks | AWS | Insurance & Financial Services
Gurugram, India | +91 8770794667 | connectwithrajeevr@gmail.com | linkedin.com/in/lethalrajeev

SUMMARY

Senior Data Engineer with 7+ years of experience building high-performance data pipelines for Insurance and Financial Services clients including CHUBB, Nationwide Mutual, and Marsh McLennan. I have migrated 50+ Informatica mappings to PySpark on Databricks, cut pipeline runtimes by 83%, and built systems that process 1M+ records daily. I work at the intersection of cloud architecture (Azure & AWS) and real-world delivery — not just building pipelines, but making them faster, cheaper, and more resilient.

TECHNICAL SKILLS

Languages: PySpark, Python, SQL
Big Data & Lakehouse: Apache Spark, Databricks Lakehouse, Delta Lake, Unity Catalog
Cloud – Azure: Azure Data Factory (ADF), Azure Synapse Analytics, ADLS Gen2, Azure Monitor, KQL
Cloud – AWS: S3, EC2, IAM, Lambda, CloudWatch, CloudFormation, SNS
ETL & Migration: Informatica, Databricks ETL Pipelines, On-prem to Cloud Migration
AI / GenAI: Generative AI, RAG, LangChain — applied to ETL automation and data tooling
CI/CD & DevOps: GitHub, Jenkins, Git, CloudFormation (IaaC)
Collaboration: Jira, ServiceNow, Agile / Sprint delivery

PROFESSIONAL EXPERIENCE

Accenture (Client: Marsh McLennan) (Client: Marsh McLennan) Jan 2026 – Present
Senior Data Engineer *Gurugram, India*

- Delivered end-to-end deployment of 2 integrations within 3 months, achieving a 96% successful build rate — ensuring on-time delivery with high reliability.
- Continued development and optimisation of Databricks ETL pipelines in a large-scale BFSI environment, maintaining high availability and data quality standards.

Accenture (Client: CHUBB Insurance) (Client: CHUBB Insurance)Oct 2024 – Dec 2025
Senior Data EngineerGurugram, India

- Cut pipeline execution time from 3 hours to under 20 minutes by re-architecting ADF pipelines with advanced performance tuning on Azure Databricks.
- Migrated 50+ Informatica ETL mappings to PySpark on Databricks — eliminating legacy on-prem dependencies and reducing migration timelines by weeks.
- Achieved 99.9% pipeline uptime across critical insurance data workflows by building fault-tolerant ADF pipelines with retry logic and monitoring.
- Lowered cloud processing costs by 35% through custom SQL optimisation within ADF, reducing ingestion load and compute overhead.
- Designed end-to-end data integration architecture connecting 10+ source systems to a unified Databricks Lakehouse — enabling seamless cross-domain analytics.
- Enforced enterprise-grade data governance across all migration pipelines — zero compliance breaches across security, audit, and data quality checkpoints.

Cognizant (Client: Nationwide Mutual Insurance) (Client: Nationwide Mutual Insurance)Feb 2022 – Oct 2024
Data EngineerKolkata, India

- Processed 1M+ insurance records daily through PySpark pipelines on Databricks — with consistent schema alignment and 99% data reliability.
- Reduced PySpark ETL runtime by 40% through partition tuning, query optimisation, and replacing redundant transformations with reusable notebook functions.
- Built a fully automated S3-to-Databricks ingestion pipeline using AWS Lambda triggers and CloudFormation — removing manual handoffs and cutting ingestion lag by 60%.
- Reduced Databricks alert response time by 70% by implementing SNS-based notification workflows integrated with SMTP and AWS CloudWatch monitoring.
- Administered Unity Catalog and Databricks workspace governance — managing role-based data access for 50+ end-users across computational and data domains.
- Accelerated deployment cycles by automating S3 resource provisioning via CloudFormation templates and AWS Code Pipeline — eliminating manual infra setup.

Mindtree Limited (Client: Microsoft) (Client: Microsoft)Nov 2018 – Feb 2022
Power BI Engineer (SRE)Hyderabad, India

- Maintained 99.9% Power BI service availability by resolving SEV1 and SEV2 incidents — reducing mean time to resolution across global outages.
- Cut report refresh time by 50% through SQL query and stored procedure optimisation — directly improving downstream analytics performance.
- Automated operational monitoring for Power BI workloads using Azure automation scripts, Power BI APIs, and KQL dashboards — reducing manual ops overhead by 30%.
- Orchestrated data ingestion pipelines via Azure Monitor, ADF, and Azure SQL Database — supporting real-time dashboards used by Microsoft teams globally.